

September 2, 2026

Dear Guest Editors,

Please consider the enclosed manuscript, "Who Authorizes Software Behavior? Governing the AI-Native SDLC," for the special issue "The Future of Software Engineering in an AI-Native World."

Software is approved at many points—code review at merge, security review before release, change approval before deployment, sign-off before configuration takes effect. Each approval covers one step; none carries forward to the next. The chain holds together because people hold its context together, which worked while software producers were few, known, and governed by shared unwritten rules. AI agents break all three conditions. The resulting system records what happened, but not necessarily what authorized a particular behavior to run.

The paper proposes a different arrangement: determine permitted behavior during construction, seal that determination into the software's state, and let execution realize it without being able to enlarge it. Proposing software and authorizing software become separate acts: an agent may produce code, but cannot thereby authorize it. The contribution is architectural rather than another technique for coding with agents.

Three forms of support can be judged separately. The architecture is an argument from reasoning. A reference implementation demonstrates feasibility across seven governed domains under one governance surface. An open specification captures the model as an implementation-independent specification, and a separate realization was built from it without access to the reference implementation.

PGC emerged from decades of building and operating large computing systems, The work is original, is not under consideration elsewhere, and I am its sole author.

Thank you for considering the manuscript.

Sincerely,

Bhash Ganti
Independent Researcher
omnibachi.org